

LoanTap Case Study

Introduction

LoanTap is at the forefront of offering tailored financial solutions to millennials. Their innovative approach seeks to harness data science for refining their credit underwriting process. The focus here is the Personal Loan segment. A deep dive into the dataset can reveal patterns in borrower behaviour and creditworthiness. Analyzing this dataset can provide crucial insights into the financial behaviours, spending habits and potential risk associated with each borrower. The insights gained can optimize loan disbursal, balancing customer outreach with risk management.

What is expected

Assuming you are a data scientist at LoanTap, you are tasked with analyzing the dataset to determine the creditworthiness of potential borrowers. Your ultimate objective is to build a logistic regression model, evaluate its performance, and provide actionable insights for the underwriting process.

1. Data

The analysis was done on the data located at -

https://drive.google.com/file/d/1ZPYj7CZCfxntE8p2Lze_4QO4MyEOy6_d/view?usp=sharing

2. Libraries

Below are the libraries required

```
In [1]: # libraries to analyze data
import numpy as np
import pandas as pd

# libraries to visualize data
import matplotlib.pyplot as plt
import seaborn as sns

import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor

from sklearn.impute import KNNImputer
from sklearn.preprocessing import OneHotEncoder, MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, ConfusionMatrixDisp
from sklearn.pipeline import make_pipeline

from imblearn.over_sampling import SMOTE
```

3. Data Loading

Loading the data into Pandas dataframe for easily handling of data

```
In [2]: # read the file into a pandas dataframe
customer_df = pd.read_csv('LoanTapData.csv')
df = customer_df
# look at the datatypes of the columns
print('*****')
print(df.info())
print('*****\n')
print('*****')
print(f'Shape of the dataset is {df.shape}')
print('*****\n')
print('*****')
print(f'Number of nan/null values in each column: \n{df.isna().sum()}')
print('*****\n')
print('*****')
print(f'Number of unique values in each column: \n{df.nunique()}')
print('*****\n')
print('*****')
print(f'Duplicate entries: \n{df.duplicated().value_counts()}')
print('*****')
```

```
*****
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                            396030 non-null  float64
1   term                                396030 non-null  object
2   int_rate                            396030 non-null  float64
3   installment                         396030 non-null  float64
4   grade                               396030 non-null  object
5   sub_grade                           396030 non-null  object
6   emp_title                           373103 non-null  object
7   emp_length                          377729 non-null  object
8   home_ownership                      396030 non-null  object
9   annual_inc                          396030 non-null  float64
10  verification_status                 396030 non-null  object
11  issue_d                             396030 non-null  object
12  loan_status                         396030 non-null  object
13  purpose                             396030 non-null  object
14  title                               394274 non-null  object
15  dti                                 396030 non-null  float64
16  earliest_cr_line                    396030 non-null  object
17  open_acc                            396030 non-null  float64
18  pub_rec                             396030 non-null  float64
19  revol_bal                           396030 non-null  float64
20  revol_util                          395754 non-null  float64
21  total_acc                           396030 non-null  float64
22  initial_list_status                 396030 non-null  object
23  application_type                    396030 non-null  object
24  mort_acc                            358235 non-null  float64
25  pub_rec_bankruptcies                395495 non-null  float64
26  address                             396030 non-null  object
dtypes: float64(12), object(15)
memory usage: 81.6+ MB
None
*****
```

```
*****
Shape of the dataset is (396030, 27)
*****
```

```
*****
```

Number of nan/null values in each column:

loan_amnt	0
term	0
int_rate	0
installment	0
grade	0
sub_grade	0
emp_title	22927
emp_length	18301
home_ownership	0
annual_inc	0
verification_status	0
issue_d	0
loan_status	0
purpose	0
title	1756
dti	0
earliest_cr_line	0
open_acc	0
pub_rec	0
revol_bal	0
revol_util	276
total_acc	0
initial_list_status	0
application_type	0
mort_acc	37795
pub_rec_bankruptcies	535
address	0

dtype: int64

Number of unique values in each column:

loan_amnt	1397
term	2
int_rate	566
installment	55706
grade	7
sub_grade	35
emp_title	173105
emp_length	11
home_ownership	6
annual_inc	27197
verification_status	3
issue_d	115
loan_status	2
purpose	14
title	48816
dti	4262
earliest_cr_line	684
open_acc	61
pub_rec	20
revol_bal	55622
revol_util	1226
total_acc	118
initial_list_status	2
application_type	3
mort_acc	33
pub_rec_bankruptcies	9
address	393700

dtype: int64

Duplicate entries:

False 396030

Name: count, dtype: int64

```
In [3]: # look at the top 5 rows
df.head()
```

Out[3]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_ownership	annual_inc
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	10+ years	RENT	11700
1	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	4 years	MORTGAGE	6500
2	15600.0	36 months	10.49	506.97	B	B3	Statistician	< 1 year	RENT	43000
3	7200.0	36 months	6.49	220.65	A	A2	Client Advocate	6 years	RENT	54000
4	24375.0	60 months	17.27	609.33	C	C5	Destiny Management Inc.	9 years	MORTGAGE	55000

5 rows × 27 columns

Not all columns are visible, so looking at groups of columns

```
In [4]: df[df.columns[10:20]]
```

Out[4]:

	verification_status	issue_d	loan_status	purpose	title	dti	earliest_cr_line	open_acc
0	Not Verified	Jan-2015	Fully Paid	vacation	Vacation	26.24	Jun-1990	16.0
1	Not Verified	Jan-2015	Fully Paid	debt_consolidation	Debt consolidation	22.05	Jul-2004	17.0
2	Source Verified	Jan-2015	Fully Paid	credit_card	Credit card refinancing	12.79	Aug-2007	13.0
3	Not Verified	Nov-2014	Fully Paid	credit_card	Credit card refinancing	2.60	Sep-2006	6.0
4	Verified	Apr-2013	Charged Off	credit_card	Credit Card Refinance	33.95	Mar-1999	13.0
...
396025	Source Verified	Oct-2015	Fully Paid	debt_consolidation	Debt consolidation	15.63	Nov-2004	6.0
396026	Source Verified	Feb-2015	Fully Paid	debt_consolidation	Debt consolidation	21.45	Feb-2006	6.0
396027	Verified	Oct-2013	Fully Paid	debt_consolidation	pay off credit cards	17.56	Mar-1997	15.0
396028	Verified	Aug-2012	Fully Paid	debt_consolidation	Loanforpayoff	15.88	Nov-1990	9.0
396029	Verified	Jun-2010	Fully Paid	debt_consolidation	Toxic Debt Payoff	8.32	Sep-1998	3.0

396030 rows × 10 columns

```
In [5]: df[df.columns[20:]]
```

Out[5]:

	revol_util	total_acc	initial_list_status	application_type	mort_acc	pub_rec_bankruptcies	
0	41.8	25.0	w	INDIVIDUAL	0.0	0.0	0174 Gateway\r\nMendocino
1	53.3	27.0	f	INDIVIDUAL	3.0	0.0	1076 Carney I 347\r\nLoganm
2	92.2	26.0	f	INDIVIDUAL	0.0	0.0	87025 Mark I 269\r\nNew Sab
3	21.5	13.0	f	INDIVIDUAL	0.0	0.0	Ford\r\nDelacruz
4	69.8	43.0	f	INDIVIDUAL	1.0	0.0	6 Roads\r\nGreggs
...	
396025	34.3	23.0	w	INDIVIDUAL	0.0	0.0	12951 Crossing\r\nJohn D
396026	95.7	8.0	f	INDIVIDUAL	1.0	0.0	0114 Fowler Fie 028\r\nRachelbor
396027	66.9	23.0	f	INDIVIDUAL	0.0	0.0	953 Matthew Poir 414\r\nReedfor
396028	53.8	20.0	f	INDIVIDUAL	5.0	0.0	7843 Blake Free 229\r\nNew Mich
396029	91.3	19.0	f	INDIVIDUAL	NaN	0.0	787 Causeway\r\nBria A

396030 rows × 7 columns

```
In [6]: df.describe()
```

Out[6]:

	loan_amnt	int_rate	installment	annual_inc	dti	open_acc	pub_rec
count	396030.000000	396030.000000	396030.000000	3.960300e+05	396030.000000	396030.000000	396030.000000
mean	14113.888089	13.639400	431.849698	7.420318e+04	17.379514	11.311153	0.178191
std	8357.441341	4.472157	250.727790	6.163762e+04	18.019092	5.137649	0.530671
min	500.000000	5.320000	16.080000	0.000000e+00	0.000000	0.000000	0.000000
25%	8000.000000	10.490000	250.330000	4.500000e+04	11.280000	8.000000	0.000000
50%	12000.000000	13.330000	375.430000	6.400000e+04	16.910000	10.000000	0.000000
75%	20000.000000	16.490000	567.300000	9.000000e+04	22.980000	14.000000	0.000000
max	40000.000000	30.990000	1533.810000	8.706582e+06	9999.000000	90.000000	86.000000

Insight

- There are **396030** entries with 27 columns
- There are **22927** null/missing values in **emp_title**, **18301** in **emp_length**, **1756** in **title**, **276** in **revol_util**, **37795** in **mort_acc** and **535** in **pub_rec_bankruptcies**
- There are no **duplicates**
- Based on the number of unique values in each column and their datatype, the columns **term**, **grade**, **emp_length**, **home_ownership**, **verification_status**, **loan_status**, **purpose**, **initial_list_status** and **application_type** can be converted to **categorical** datatype
- **issue_d** and **earliest_cr_line** has date values so can be converted to **datetime**
- The column **sub_grade** can be dropped as its info is already captured in column **grade**
- The columns **term** and **emp_length** have space, special characters, alphanumeric mix.
- Replace "36 months" with "short" and "60 months" with "long" in **term** column
- Remove "year/years", replacing "< 1" with "0" and replacing "10+" with "10" in **emp_length** column. Then convert **emp_length** column to **categorical** datatype
- Extract zip code from **address** column and drop address column

```
In [7]: # Convert to category
categorical_columns = ['term', 'grade', 'emp_length', 'home_ownership', 'verification_status']
df[categorical_columns] = df[categorical_columns].astype('category')

# Convert to datetime
df['issue_d'] = pd.to_datetime(df['issue_d'], format='%b-%Y')
df['earliest_cr_line'] = pd.to_datetime(df['earliest_cr_line'], format='%b-%Y')

# Drop "grade" column
df.drop(columns=['sub_grade'], inplace=True)

# Rename the values in 'term' column
df['term'].replace({' 36 months' : 'short', ' 60 months' : 'long'}, inplace=True)

# Rename employee_length column values
df['emp_length'] = df['emp_length'].replace({'< 1 year': '0 year'})
df['emp_length'] = df['emp_length'].str.replace(r'\D', '', regex=True)
df['emp_length'] = df['emp_length'].astype('category')

# Extract zip code from address
df['zip_code'] = df['address'].str[-5:].str.split().str[0].astype('category')
df.drop(columns='address', inplace=True)
print(f'Number of nan/null values in zip_code column: \n{df.zip_code.isna().sum()}')

df.info()
```

Number of nan/null values in zip_code column:

0

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 396030 entries, 0 to 396029

Data columns (total 26 columns):

#	Column	Non-Null Count	Dtype
0	loan_amnt	396030 non-null	float64
1	term	396030 non-null	category
2	int_rate	396030 non-null	float64
3	installment	396030 non-null	float64
4	grade	396030 non-null	category
5	emp_title	373103 non-null	object
6	emp_length	377729 non-null	category
7	home_ownership	396030 non-null	category
8	annual_inc	396030 non-null	float64
9	verification_status	396030 non-null	category

```

10  issue_d                396030 non-null  datetime64[ns]
11  loan_status            396030 non-null  category
12  purpose                396030 non-null  category
13  title                  394274 non-null  object
14  dti                    396030 non-null  float64
15  earliest_cr_line       396030 non-null  datetime64[ns]
16  open_acc               396030 non-null  float64
17  pub_rec                396030 non-null  float64
18  revol_bal              396030 non-null  float64
19  revol_util              395754 non-null  float64
20  total_acc               396030 non-null  float64
21  initial_list_status     396030 non-null  category
22  application_type        396030 non-null  category
23  mort_acc                358235 non-null  float64
24  pub_rec_bankruptcies    395495 non-null  float64
25  zip_code                396030 non-null  category
dtypes: category(10), datetime64[ns](2), float64(12), object(2)
memory usage: 52.1+ MB

```

4. Exploratory Data Analysis

4.1. Handling null values

emp_title has 22927, **emp_length** has 18301, **title** has 1756, **revol_util** has 276, **mort_acc** has 37795 and **pub_rec_bankruptcies** has 535 null values

- Columns **emp_title** and **title** can be **dropped** as they would not have an effect on the loan approval
- Null values in **revol_util** and **pub_rec_bankruptcies** are small in number and hence can be dropped

```

In [8]: df.drop(columns=['emp_title', 'title'], inplace=True)
df.dropna(subset = ['revol_util', 'pub_rec_bankruptcies'], inplace=True)



```

- Let us check the distribution of remaining features before deciding on how to handle the null values

```

In [9]: fig, axs = plt.subplots(1,2, figsize=(10,4))
sns.histplot(ax = axs[0], data=df, x = 'emp_length')
sns.histplot(ax = axs[1], data=df, x = 'mort_acc')
fig.tight_layout()
plt.show()

```



- **mort_acc** is the number of mortgage accounts out of the **total_acc**. So I will replace the null values with "mode" of **mort_acc** for different **total_acc**
- For **emp_length**, I will use knn imputer to fill the missing data

```
In [10]: mode_mort_acc_df = df.groupby('total_acc')['mort_acc'].agg(lambda x: pd.Series.mode(x)[0])

def fill_mort_acc(total_acc, mort_acc):
    if np.isnan(mort_acc):
        return mode_mort_acc_df[total_acc]
    else:
        return mort_acc

df['mort_acc'] = df.apply(lambda x: fill_mort_acc(x['total_acc'], x['mort_acc']), axis=1)
```

```
In [11]: imputer = KNNImputer(n_neighbors=5)
df['emp_length'] = imputer.fit_transform(df[['emp_length']])
df['emp_length'] = df['emp_length'].astype('category')
```

```
In [12]: df.isna().sum()
```

```
Out[12]: loan_amnt      0
term      0
int_rate   0
installment 0
grade      0
emp_length 0
home_ownership 0
annual_inc 0
verification_status 0
issue_d     0
loan_status 0
purpose     0
dti         0
earliest_cr_line 0
open_acc    0
pub_rec     0
revol_bal   0
revol_util  0
total_acc   0
initial_list_status 0
application_type 0
mort_acc    0
pub_rec_bankruptcies 0
zip_code     0
dtype: int64
```

There are no null values now

4.2. Detecting outliers

4.2.1. Outliers for every continuous variable

```
In [13]: # helper function to detect outliers using IQR method
def detectOutliers_iqr(df):
    q1 = df.quantile(0.25)
    q3 = df.quantile(0.75)
    iqr = q3 - q1
    lower_outliers = df[df < (q1 - 1.5 * iqr)]
    higher_outliers = df[df > (q3 + 1.5 * iqr)]
    return lower_outliers, higher_outliers
```




```
In [14]: # helper function to detect outliers using standard deviation method
def detectOutliers_std(df):
    mean = df.mean()
    std = df.std()
    upper_limit = mean+(3*std)
    lower_limit = mean-(3*std)
    lower_outliers = df[df<lower_limit]
    higher_outliers = df[df>upper_limit]
    return lower_outliers, higher_outliers
```

```
In [15]: numerical_columns = df.select_dtypes(include=np.number).columns
column_outlier_dictionary = {}
for column in numerical_columns:
    lower_outliers, higher_outliers = detectOutliers_iqr(df[column])
    column_outlier_dictionary[column] = [lower_outliers, higher_outliers]
    #print('*'*50)
    #print(f'Outliers of \'{column}\'' column are:')
    #print("Lower outliers:\n", lower_outliers)
    #print("Higher outliers:\n", higher_outliers)
    #print('*'*50, end="\n")
```

```
In [16]: for key, value in column_outlier_dictionary.items():
    print(f'The column \'{key}\'' has {len(value[0]) + len(value[1])} outliers')
```

```
The column 'loan_amnt' has 190 outliers
The column 'int_rate' has 3144 outliers
The column 'installment' has 11114 outliers
The column 'annual_inc' has 16649 outliers
The column 'dti' has 275 outliers
The column 'open_acc' has 10297 outliers
The column 'pub_rec' has 57730 outliers
The column 'revol_bal' has 21205 outliers
The column 'revol_util' has 12 outliers
The column 'total_acc' has 8491 outliers
The column 'mort_acc' has 6837 outliers
The column 'pub_rec_bankruptcies' has 45111 outliers
```

```
In [17]: num_cols = 6
num_rows = int(np.ceil(len(numerical_columns)/num_cols))
fig, axs = plt.subplots(num_rows, num_cols, figsize=(10,15))
for idx in range(len(numerical_columns)):
    ax = plt.subplot(num_rows, num_cols, idx+1)
    sns.boxplot(ax = ax, data=df, y = numerical_columns[idx])
plt.tight_layout()
plt.show()
```



```
In [18]: numerical_columns = df.select_dtypes(include=np.number).columns
numerical_columns = list(numerical_columns)
numerical_columns.remove('pub_rec')
numerical_columns.remove('pub_rec_bankruptcies')
numerical_columns = pd.core.indexes.base.Index(numerical_columns)
column_outlier_dictionary = {}
for column in numerical_columns:
    lower_outliers, higher_outliers = detectOutliers_std(df[column])
    column_outlier_dictionary[column] = [lower_outliers, higher_outliers]
    #print('*'*50)
    #print(f'Outliers of \'{column}\'' column are:')
    #print("Lower outliers:\n", lower_outliers)
    #print("Higher outliers:\n", higher_outliers)
    #print('*'*50, end="\n")
```

```
In [19]: for key, value in column_outlier_dictionary.items():
        print(f'The column \'{key}\'' has {len(value[0]) + len(value[1])} outliers')
```

```
The column 'loan_amnt' has 184 outliers
The column 'int_rate' has 754 outliers
The column 'installment' has 5042 outliers
The column 'annual_inc' has 3190 outliers
The column 'dti' has 12 outliers
The column 'open_acc' has 4873 outliers
The column 'revol_bal' has 4771 outliers
The column 'revol_util' has 16 outliers
The column 'total_acc' has 3396 outliers
The column 'mort_acc' has 6837 outliers
```

4.2.2. Remove the outliers

- Based on the boxplot, the number of outliers using IQR method and standard deviation method, I will remove the outliers using the standard deviation method except for columns **pub_rec** and **pub_rec_bankruptcies** which will be removed based on manual check.

```
In [20]: remove_outliers = True
if True == remove_outliers:
    master_index = pd.core.indexes.base.Index([])
    for key, value in column_outlier_dictionary.items():
        lower_outliers = value[0]
        higher_outliers = value[1]
        master_index = master_index.union(lower_outliers.index).union(higher_outliers.in
df.drop(master_index, inplace=True)
else:
    print('Not removing any outliers')
```

```
In [21]: df['pub_rec'].value_counts()
```

```
Out[21]: pub_rec
0.0      315552
1.0      47129
2.0       5107
3.0      1424
4.0       481
5.0       218
6.0       108
7.0        47
8.0        31
10.0       11
9.0        10
11.0        6
13.0        4
```


```
12.0      4
19.0      2
40.0      1
17.0      1
86.0      1
24.0      1
15.0      1
Name: count, dtype: int64
```

```
In [22]: df['pub_rec_bankruptcies'].value_counts()
```

```
Out[22]: pub_rec_bankruptcies
0.0      327200
1.0      40774
2.0       1716
3.0        332
4.0         75
5.0         30
6.0          6
7.0          4
8.0          2
Name: count, dtype: int64
```

```
In [23]: df = df[df['pub_rec'] < 11]
df = df[df['pub_rec_bankruptcies'] < 6]
```

```
In [24]: numerical_columns = df.select_dtypes(include=np.number).columns
num_cols = 6
num_rows = int(np.ceil(len(numerical_columns)/num_cols))
fig, axs = plt.subplots(num_rows, num_cols, figsize=(10,15))
for idx in range(len(numerical_columns)):
    ax = plt.subplot(num_rows, num_cols, idx+1)
    sns.boxplot(ax = ax, data=df, y = numerical_columns[idx])
plt.tight_layout()
plt.show()
```



```
In [25]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 370106 entries, 0 to 396029
Data columns (total 24 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                            370106 non-null  float64
1   term                                370106 non-null  category
2   int_rate                            370106 non-null  float64
3   installment                         370106 non-null  float64
4   grade                               370106 non-null  category
5   emp_length                          370106 non-null  category
6   home_ownership                     370106 non-null  category
7   annual_inc                         370106 non-null  float64
8   verification_status               370106 non-null  category
9   issue_d                           370106 non-null  datetime64[ns]
10  loan_status                        370106 non-null  category
11  purpose                           370106 non-null  category
12  dti                               370106 non-null  float64
13  earliest_cr_line                  370106 non-null  datetime64[ns]
14  open_acc                         370106 non-null  float64
15  pub_rec                         370106 non-null  float64
16  revol_bal                       370106 non-null  float64
17  revol_util                      370106 non-null  float64
18  total_acc                      370106 non-null  float64
19  initial_list_status              370106 non-null  category
20  application_type                370106 non-null  category
21  mort_acc                       370106 non-null  float64
22  pub_rec_bankruptcies            370106 non-null  float64
23  zip_code                        370106 non-null  category
dtypes: category(10), datetime64[ns](2), float64(12)
memory usage: 45.9 MB
```

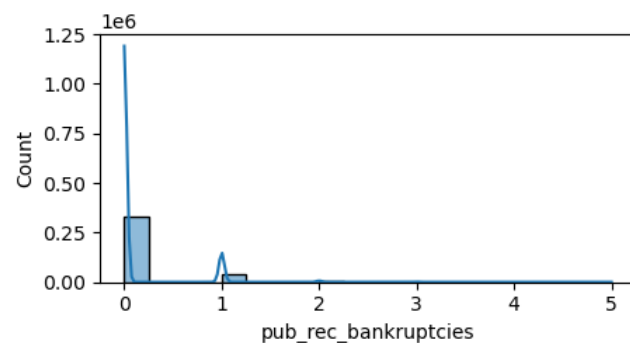
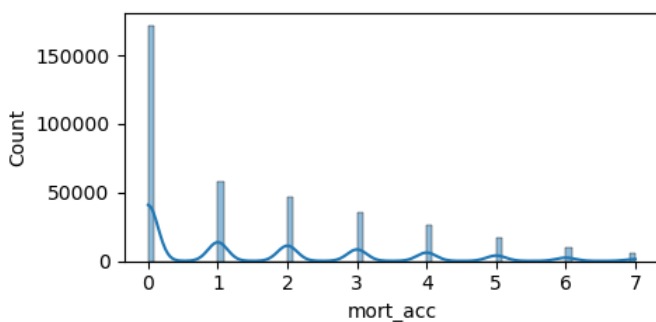

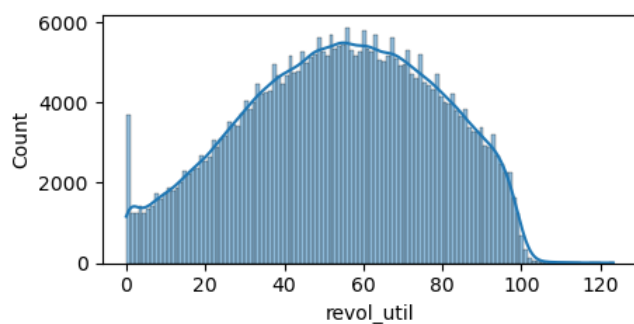
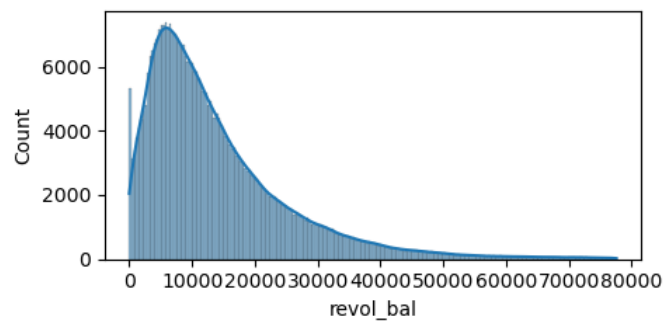

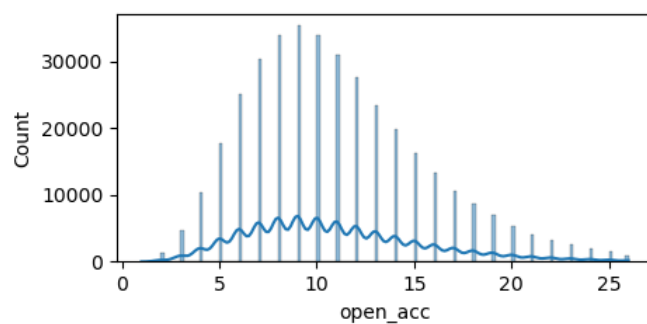





Insight

- The number of columns reduced to **370106** from the original 396030 rows

4.3. Univariate analysis

4.3.1. Numerical Variables

```
In [26]: num_cols = 2
num_rows = int(np.ceil(len(numerical_columns)/num_cols))
fig, axs = plt.subplots(num_rows, num_cols, figsize=(10,15))
for idx in range(len(numerical_columns)):
    ax = plt.subplot(num_rows, num_cols, idx+1)
    sns.histplot(ax = ax, data=df, x = numerical_columns[idx], kde=True)
plt.tight_layout()
plt.show()
```

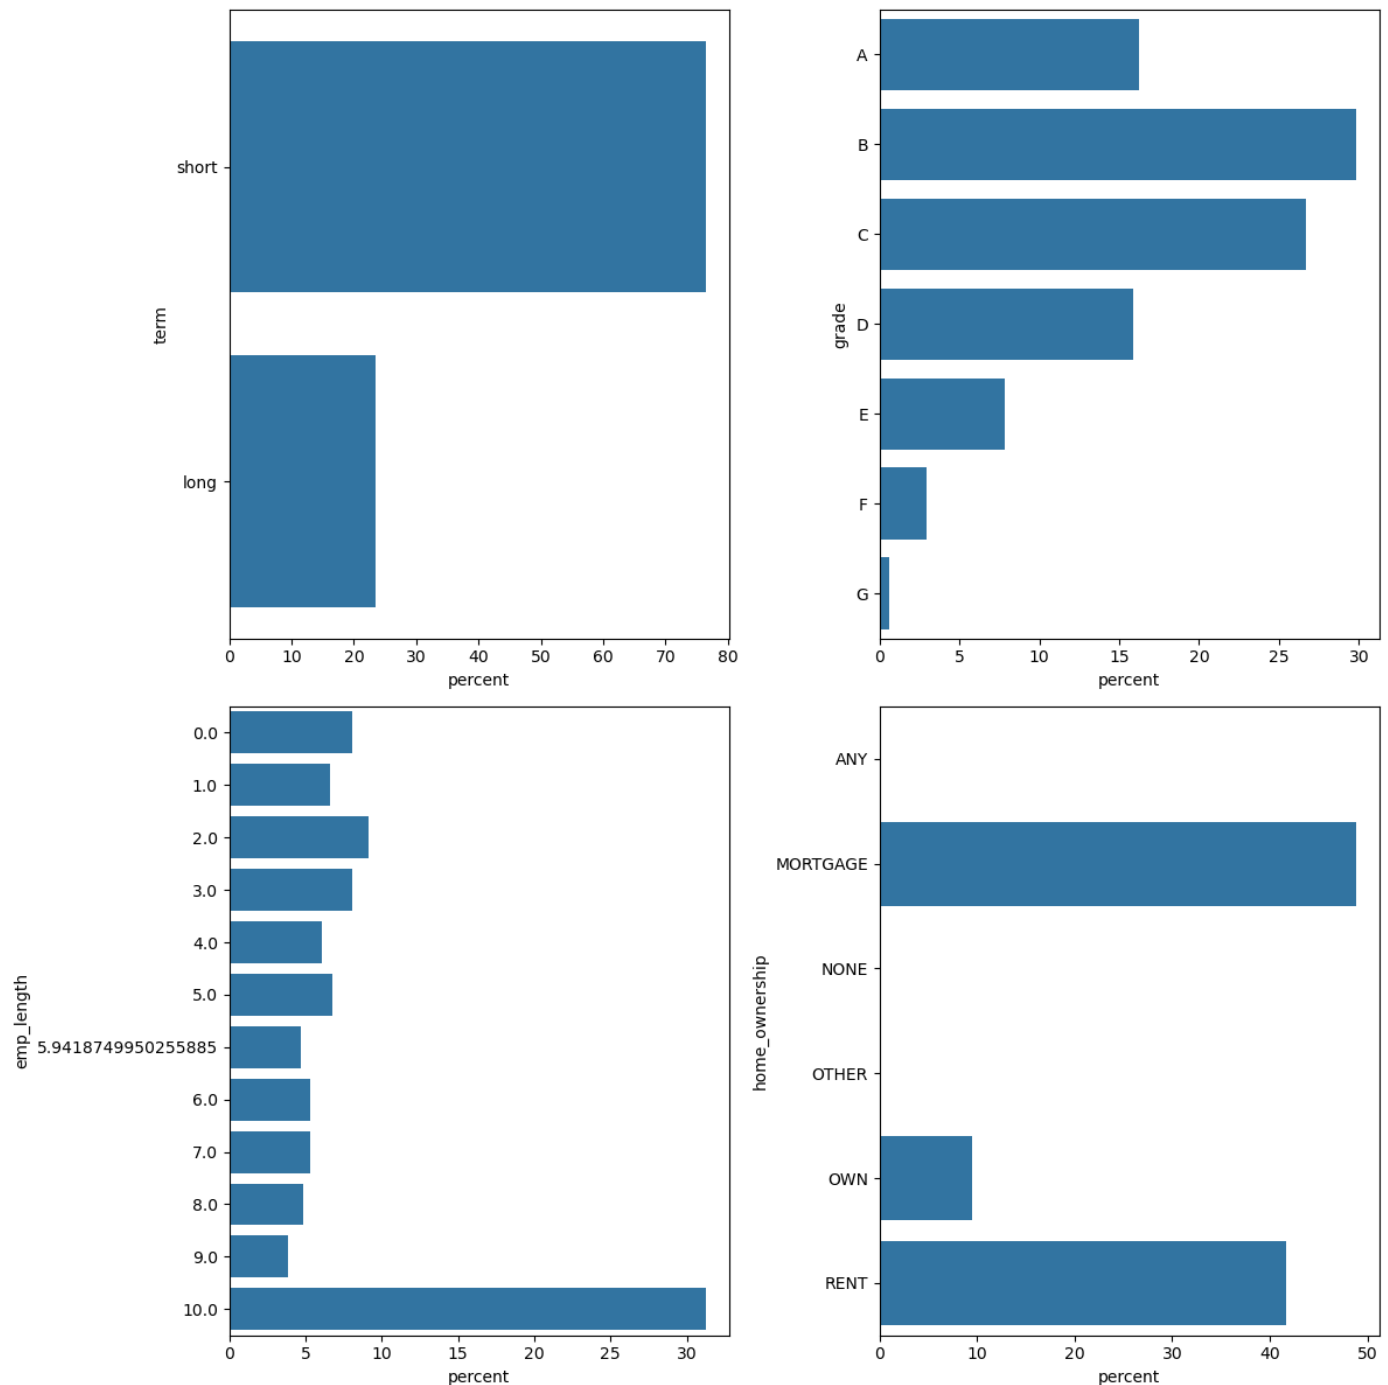


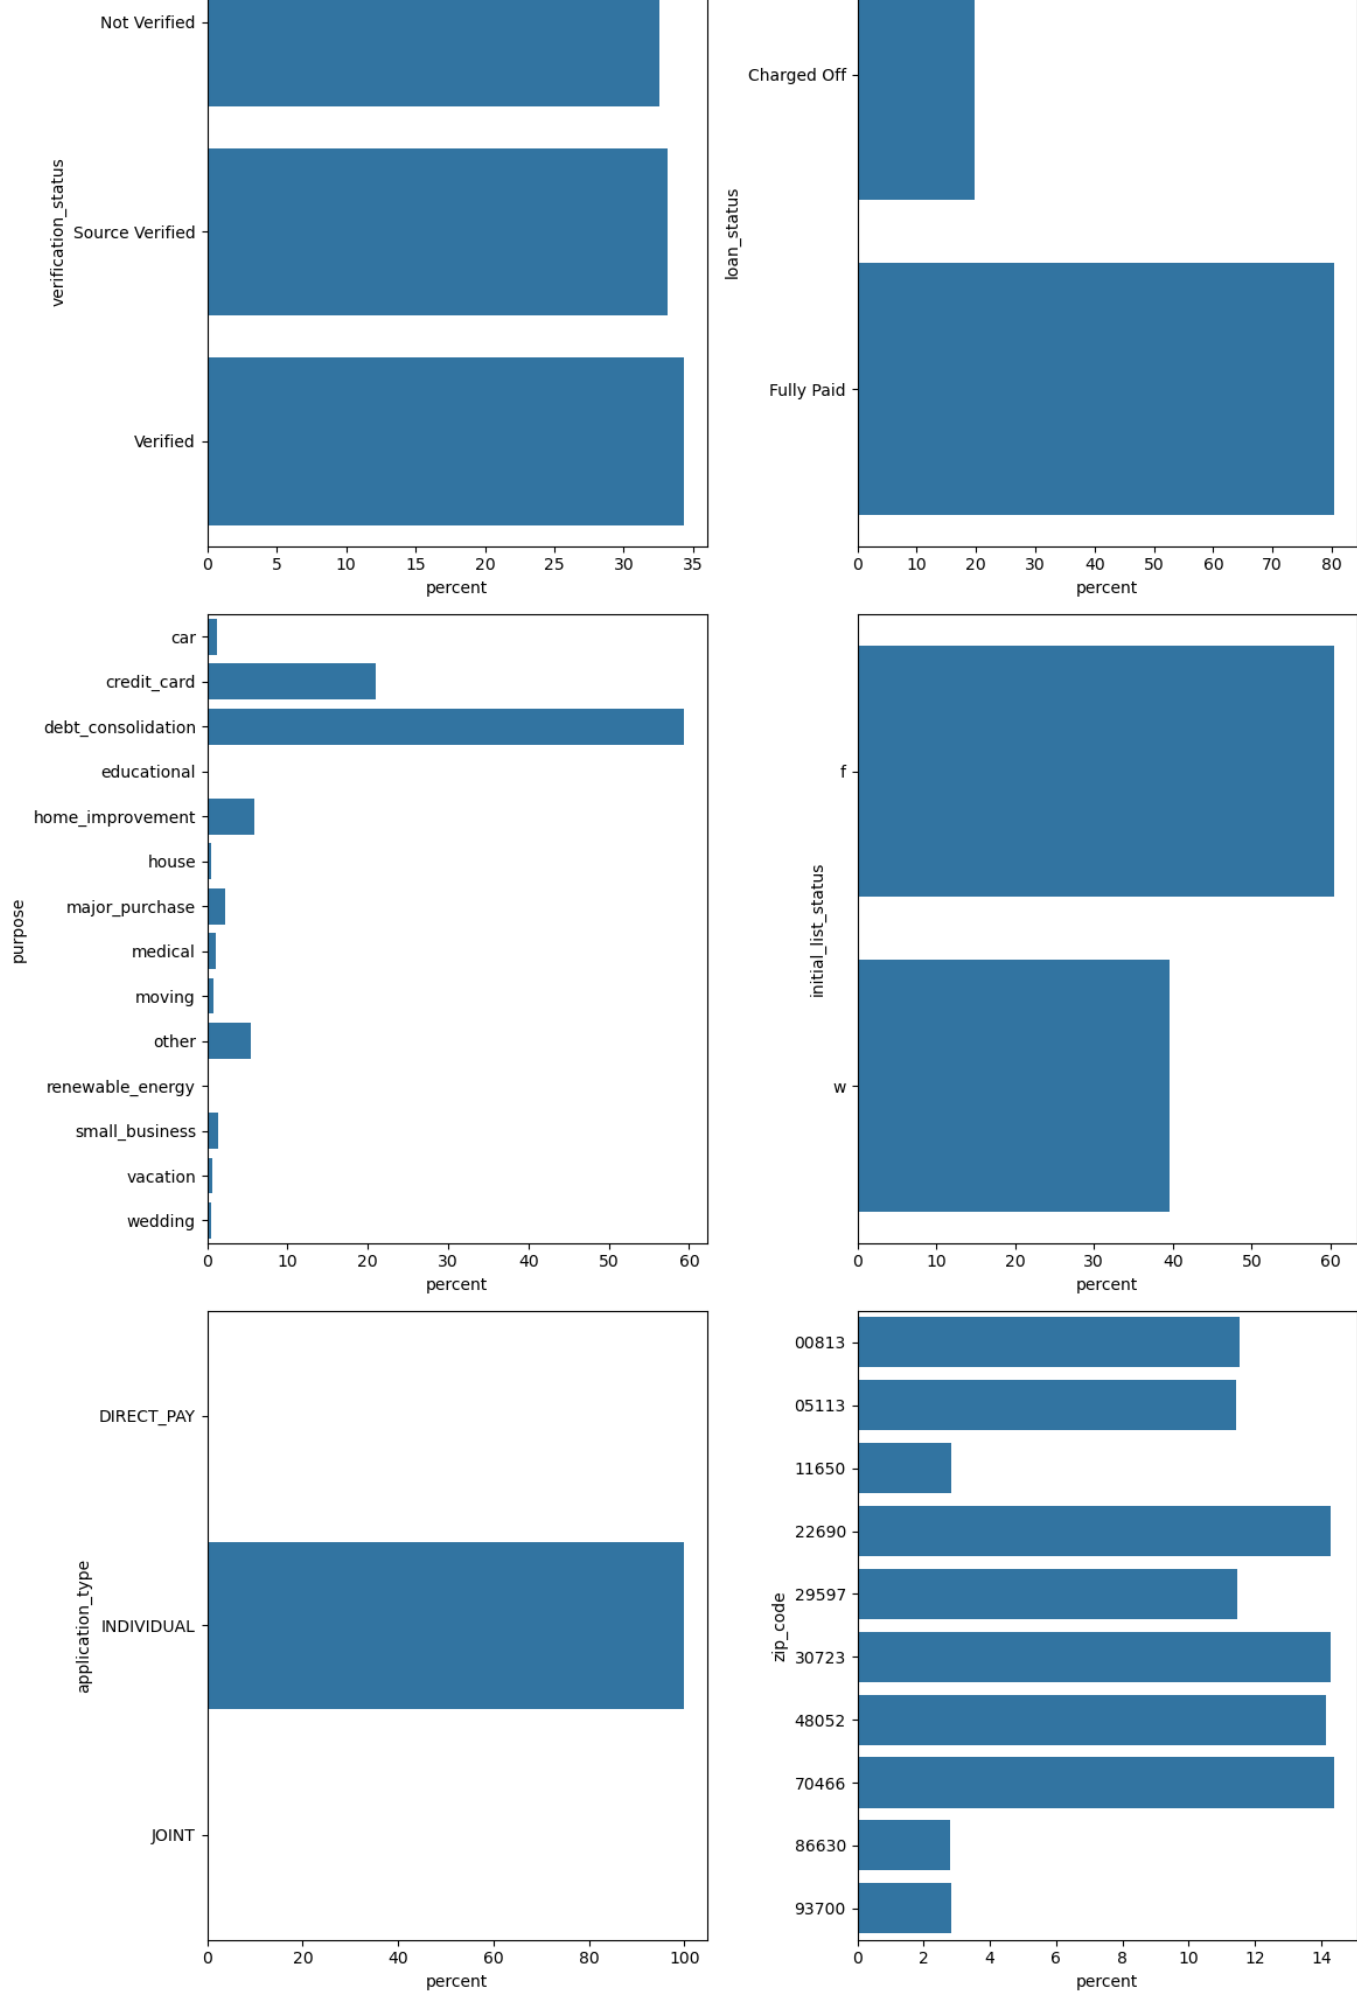
Insight

- Most of the features are right skewed except for **revol_util** which is slightly left skewed

4.3.2. Categorical Variables

```
In [27]: categorical_columns = df.select_dtypes(include='category').columns
num_cols = 2
num_rows = int(np.ceil(len(categorical_columns)/num_cols))
fig, axs = plt.subplots(num_rows, num_cols, figsize=(12, 30))
for idx in range(len(categorical_columns)):
    ax = plt.subplot(num_rows, num_cols, idx+1)
    sns.countplot(ax = ax, data=df, y = categorical_columns[idx], stat='percent')
plt.tight_layout()
plt.show()
```






- Most of the loan is taken for a short **term** which is for **3 years**
- Maximum loan are assigned the **grade B** followed by grade C
- Maximum loan are taken by borrowers whose employment tenure is more than **10 years**
- Maximum loan are taken by borrowers whose house is **mortgaged** followed by those whose who are in **rented** house
- Most of the loan is **fully_paid**
- Most of the loan is taken for **debt consolidation** followed by **credit card**
- Majority of the loan is taken by **individuals**

4.4. Bivariate analysis

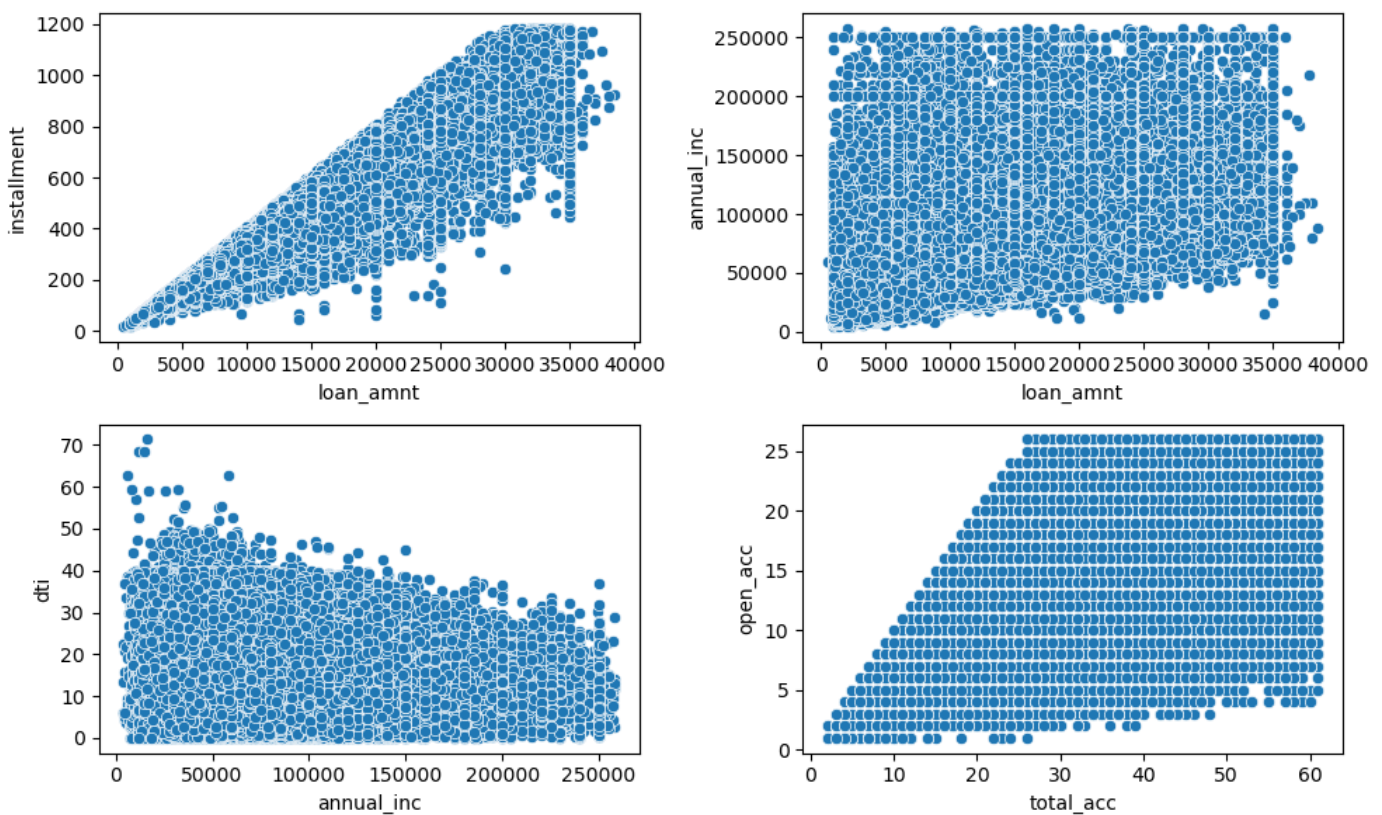
```
In [28]: sns.boxplot(data=df, x='loan_status', y='loan_amnt')
plt.show()
```



Insight

- The **median** of the **loan amount** slightly **higher** for loans which were **charged off**

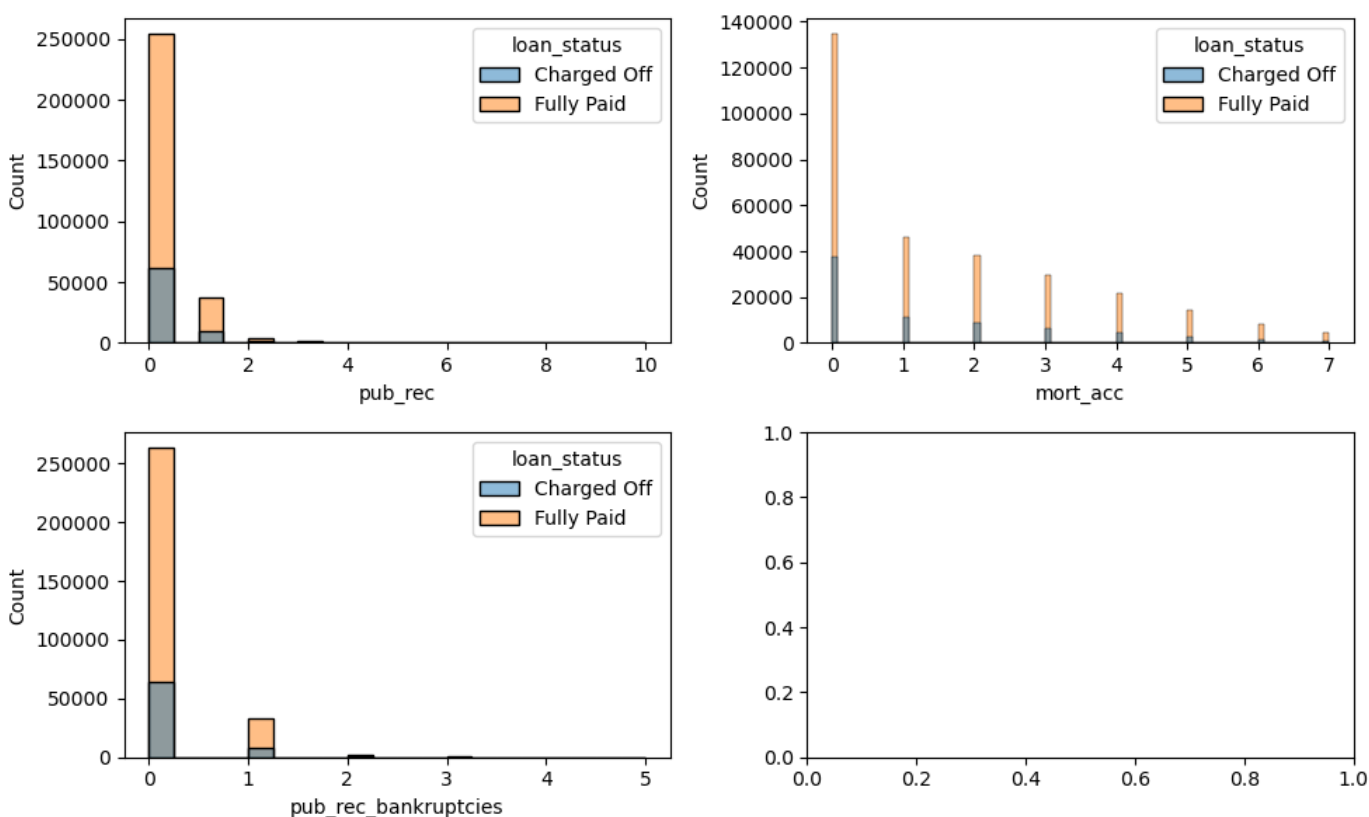
```
In [29]: fig, axs = plt.subplots(2,2,figsize=(10,6))
sns.scatterplot(ax=axs[0,0], data=df, x='loan_amnt', y='installment')
sns.scatterplot(ax=axs[0,1], data=df, x='loan_amnt', y='annual_inc')
sns.scatterplot(ax=axs[1,0], data=df, x='annual_inc', y='dti')
sns.scatterplot(ax=axs[1,1], data=df, x='total_acc', y='open_acc')
fig.tight_layout()
plt.show()
```



Insight

- It is very clear from the plot that, in general, the **installment increases** as the **loan amount** increases
- It is very obvious that people with **higher income** can afford to take **higher loan**
- As **income** increases **debt to income ratio** reduces
- Most of the borrower's accounts are active accounts

```
In [30]: fig, axs = plt.subplots(2,2,figsize=(10,6))
sns.histplot(ax=axs[0,0], data=df, x='pub_rec', hue='loan_status')
sns.histplot(ax=axs[0,1], data=df, x='mort_acc', hue='loan_status')
sns.histplot(ax=axs[1,0], data=df, x='pub_rec_bankruptcies', hue='loan_status')
fig.tight_layout()
plt.show()
```



Insight

- I will group **`pub_rec`** into 2 groups: **0** for having 0 negative records and **1** for having more than 0 negative records
- Same for **`mort_acc`** and **`pub_rec_bankruptcies`**



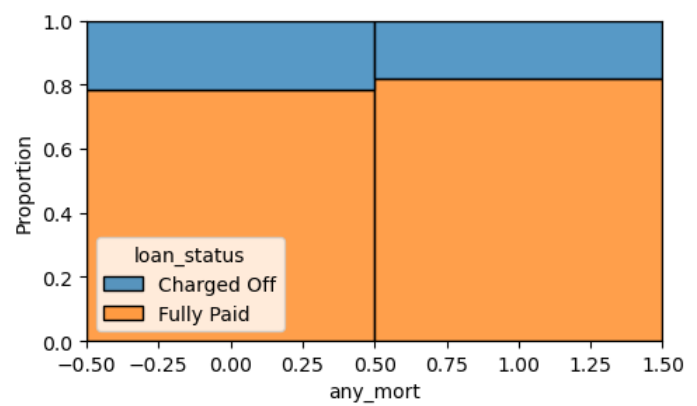
```
In [31]: group_0_list = [0.0]
pub_rec_list = list(df['pub_rec'].explode().unique())
group_1_list = list(set(pub_rec_list) - set(group_0_list))
df['any_neg_rec'] = df['pub_rec'].replace(group_0_list, 0)
df['any_neg_rec'] = df['any_neg_rec'].replace(group_1_list, 1)
df['any_neg_rec'] = df['any_neg_rec'].astype('category')

group_0_list = [0.0]
pub_rec_list = list(df['mort_acc'].explode().unique())
group_1_list = list(set(pub_rec_list) - set(group_0_list))
df['any_mort'] = df['mort_acc'].replace(group_0_list, 0)
df['any_mort'] = df['any_mort'].replace(group_1_list, 1)
df['any_mort'] = df['any_mort'].astype('category')

group_0_list = [0.0]
pub_rec_list = list(df['pub_rec_bankruptcies'].explode().unique())
group_1_list = list(set(pub_rec_list) - set(group_0_list))
df['any_bankruptcies'] = df['pub_rec_bankruptcies'].replace(group_0_list, 0)
df['any_bankruptcies'] = df['any_bankruptcies'].replace(group_1_list, 1)
df['any_bankruptcies'] = df['any_bankruptcies'].astype('category')







df.drop(columns = ['pub_rec', 'mort_acc', 'pub_rec_bankruptcies'], inplace=True)
```

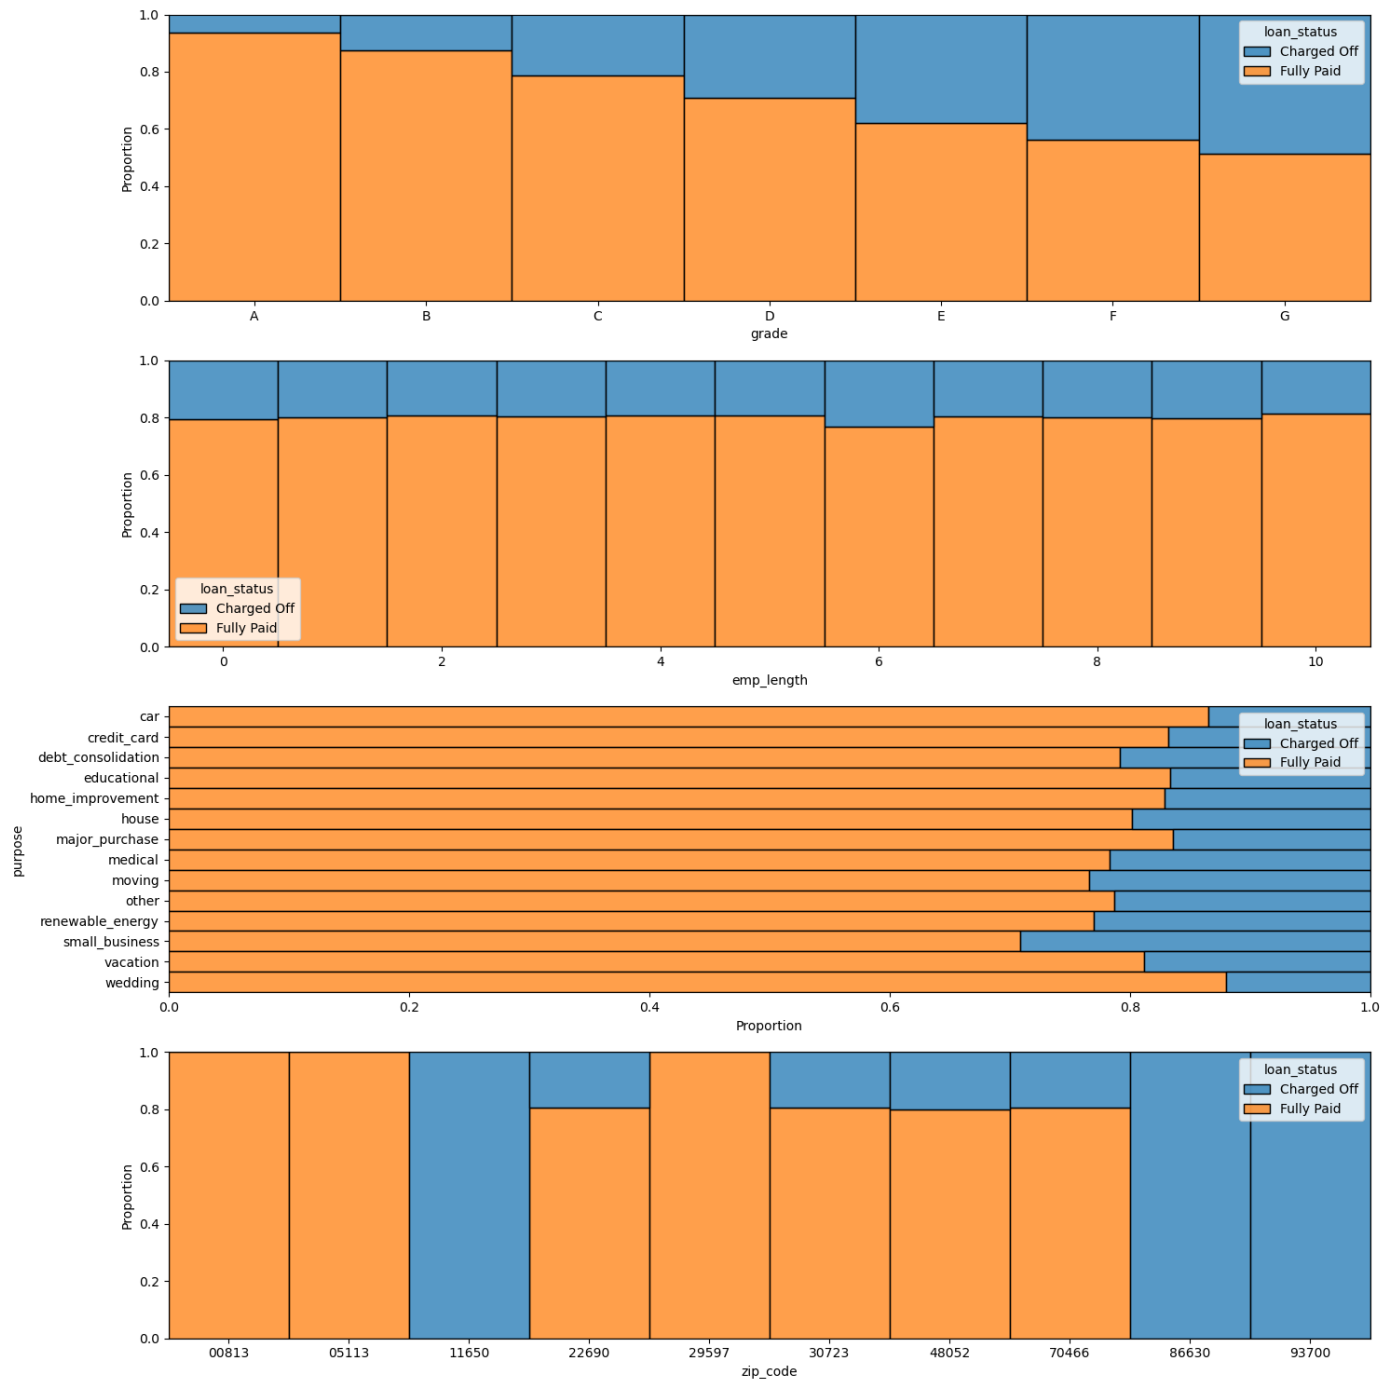
```
In [32]: fig, axs = plt.subplots(2,2,figsize=(10,6))
sns.histplot(ax=axs[0,0], data=df, x='any_neg_rec', hue='loan_status', stat="proportion")
sns.histplot(ax=axs[0,1], data=df, x='any_mort', hue='loan_status', stat="proportion", m
sns.histplot(ax=axs[1,0], data=df, x='any_bankruptcies', hue='loan_status', stat="propor
fig.tight_layout()
plt.show()
```



```
In [33]: fig, axs = plt.subplots(3,2,figsize=(12,10))
sns.histplot(ax=axs[0,0], data=df, x='term', hue='loan_status', stat="proportion", multi
sns.histplot(ax=axs[0,1], data=df, x='home_ownership', hue='loan_status', stat="proporti
sns.histplot(ax=axs[1,0], data=df, x='verification_status', hue='loan_status', stat="pro
sns.histplot(ax=axs[1,1], data=df, x='initial_list_status', hue='loan_status', stat="pro
sns.histplot(ax=axs[2,0], data=df, x='application_type', hue='loan_status', stat="propor
fig.tight_layout()
plt.show()

fig, axs = plt.subplots(4,1,figsize=(15,15))
sns.histplot(ax=axs[0], data=df, x='grade', hue='loan_status', stat="proportion", multip
sns.histplot(ax=axs[1], data=df, x='emp_length', hue='loan_status', stat="proportion", m
sns.histplot(ax=axs[2], data=df, y='purpose', hue='loan_status', stat="proportion", mult
sns.histplot(ax=axs[3], data=df, x='zip_code', hue='loan_status', stat="proportion", mul
fig.tight_layout()
plt.show()
```






Insight

- Having a **negative** or **bankruptcy** record **doesn't** seem to **impact** the loan getting paid back or not
- Loan taken for short term, i.e. **3 years** are most likely to be **fully paid** back
- Loan taken by people whose **house ownership** is of type **any** is on are **paid back 100%**.
- Surprisingly, loans which are **not verified** are more likely to be **paid back**
- **Initial list status** also **do not** seem to **impact loan status**
- Loan taken as **joint** application type are more likely to be **paid back**
- Loans with **grade A and B** are more likely to get paid back
- **Employment duration does not** seem to **impact loan status**
- Loan taken for **wedding** are more likely to be paid back
- Loan taken by people with zip code **00813 and 05113** are **fully paid** back whereas loan taken by people with zip code **11650, 86630 and 93700** are all **charged off**

In [34]: `df.drop(columns=['any_neg_rec', 'any_bankruptcies', 'initial_list_status', 'emp_length'])`

4.5. Multivariate analysis

```
In [35]: fig, ax = plt.subplots(figsize=(8,6))
sns.heatmap(df.select_dtypes(include=np.number).corr(), annot=True, linewidth=0.5, cmap=
ax.xaxis.tick_top()
plt.xticks(rotation=30, ha='left')
plt.show()
```



Insight

- **loan amount** is highly correlated with **installment**
- There is good correlation between **loan amount - annual income**, **loan amount - revol balance**, **installment - annual income**, **installment - revol balance**, **open account - total account**

5. Data Preprocessing

```
In [36]: df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 370106 entries, 0 to 396029
Data columns (total 20 columns):
 #   Column              Non-Null Count  Dtype  
---  -
 0   loan_amnt           370106 non-null float64
```



```

1 term 370106 non-null category
2 int_rate 370106 non-null float64
3 installment 370106 non-null float64
4 grade 370106 non-null category
5 home_ownership 370106 non-null category
6 annual_inc 370106 non-null float64
7 verification_status 370106 non-null category
8 issue_d 370106 non-null datetime64[ns]
9 loan_status 370106 non-null category
10 purpose 370106 non-null category
11 dti 370106 non-null float64
12 earliest_cr_line 370106 non-null datetime64[ns]
13 open_acc 370106 non-null float64
14 revol_bal 370106 non-null float64
15 revol_util 370106 non-null float64
16 total_acc 370106 non-null float64
17 application_type 370106 non-null category
18 zip_code 370106 non-null category
19 any_mort 370106 non-null category
dtypes: category(9), datetime64[ns](2), float64(9)
memory usage: 37.1 MB

```

The date features will not have an impact on the loan status, so i will drop ***issue_d*** and ***earliest_cr_line*** columns

```
In [37]: df.drop(columns=['issue_d', 'earliest_cr_line'], inplace=True)
```

5.1. Multicollinearity Check

```
In [38]: features_df = df.drop(columns=['loan_status']) # Drop target column
features_df.select_dtypes(include='category').columns
features_df = features_df.drop(columns=features_df.select_dtypes(include='category').columns)
features_df = sm.add_constant(features_df) # Adding a constant column for the intercept
vif_df = pd.DataFrame()
vif_df['Features'] = features_df.columns
vif_df['VIF'] = [variance_inflation_factor(features_df.values, idx) for idx in range(len(features_df.columns))]
vif_df['VIF'] = round(vif_df['VIF'], 2)
vif_df = vif_df.sort_values(by='VIF', ascending=False)
vif_df

```

```
Out[38]:
```

	Features	VIF
0	const	24.91
1	loan_amnt	11.58
3	installment	11.04
6	open_acc	2.00
9	total_acc	1.87
7	revol_bal	1.76
4	annual_inc	1.62
8	revol_util	1.47
5	dti	1.40
2	int_rate	1.23

- **loan amount** is highly correlated with **installment** which is also shown here by high VIF values. I will drop **installment**

```
In [39]: features_df = features_df.drop(columns=['installment'])
features_df = sm.add_constant(features_df) # Adding a constant column for the intercept
vif_df = pd.DataFrame()
vif_df['Features'] = features_df.columns
vif_df['VIF'] = [variance_inflation_factor(features_df.values, idx) for idx in range(len(features_df.columns))]
vif_df['VIF'] = round(vif_df['VIF'], 2)
vif_df = vif_df.sort_values(by='VIF', ascending=False)
vif_df
```

```
Out[39]:
```

	Features	VIF
0	const	24.76
5	open_acc	2.00
8	total_acc	1.86
6	revol_bal	1.75
3	annual_inc	1.62
1	loan_amnt	1.49
7	revol_util	1.46
4	dti	1.40
2	int_rate	1.22

Insight

- Based on the above VIF scores, I can conclude that there are no more multicollinear numerical features
- I will drop **installment** from the dataframe

```
In [40]: df.drop(columns=['installment'], inplace=True)
```

5.2. Encode categorical variables

```
In [41]: final_df = df.copy()
final_df.reset_index(inplace=True, drop=True)
```

Sepearte out target and feature columns

```
In [42]: X = final_df.drop(columns=['loan_status'])
y = final_df['loan_status']
```

Encode target variable

```
In [43]: y = y.replace({'Fully Paid': 0, 'Charged Off': 1}).astype(int)
```

Encode features with just 2 classes as 0 or 1

```
In [44]: X['term'] = X['term'].replace({'short': 0, 'long': 1}).astype('int8')
X['any_mort'] = X['any_mort'].astype('int8')
```

One-Hot-Encoding for remaining categorical features

In [45]: `X.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 370106 entries, 0 to 370105
Data columns (total 16 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   loan_amnt             370106 non-null float64
 1   term                  370106 non-null int8    
 2   int_rate              370106 non-null float64
 3   grade                 370106 non-null category
 4   home_ownership        370106 non-null category
 5   annual_inc            370106 non-null float64
 6   verification_status   370106 non-null category
 7   purpose               370106 non-null category
 8   dti                   370106 non-null float64
 9   open_acc              370106 non-null float64
10   revol_bal             370106 non-null float64
11   revol_util            370106 non-null float64
12   total_acc             370106 non-null float64
13   application_type      370106 non-null category
14   zip_code              370106 non-null category
15   any_mort              370106 non-null int8    
dtypes: category(6), float64(8), int8(2)
memory usage: 25.4 MB
```

In [46]: `categorical_columns = X.select_dtypes(include='category').columns`
`categorical_columns`

Out[46]: `Index(['grade', 'home_ownership', 'verification_status', 'purpose',
 'application_type', 'zip_code'],
 dtype='object')`

In [47]: `encoder = OneHotEncoder(sparse_output=False)`
`encoded_data = encoder.fit_transform(X[categorical_columns])`
`encoded_df = pd.DataFrame(encoded_data, columns = encoder.get_feature_names_out(categori`
`X = pd.concat([X, encoded_df], axis=1)`
`X.drop(columns = categorical_columns, inplace=True)`
`X.head()`

Out[47]:

	loan_amnt	term	int_rate	annual_inc	dti	open_acc	revol_bal	revol_util	total_acc	any_mort	...	zip_code
0	10000.0	0	11.44	117000.0	26.24	16.0	36369.0	41.8	25.0	0	...	
1	8000.0	0	11.99	65000.0	22.05	17.0	20131.0	53.3	27.0	1	...	
2	15600.0	0	10.49	43057.0	12.79	13.0	11987.0	92.2	26.0	0	...	
3	7200.0	0	6.49	54000.0	2.60	6.0	5472.0	21.5	13.0	0	...	
4	24375.0	1	17.27	55000.0	33.95	13.0	24584.0	69.8	43.0	1	...	

5 rows × 53 columns

5.3. Train-test split

In [48]: `X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)`
`X_train.shape, X_test.shape, y_train.shape, y_test.shape`

Out[48]: `((296084, 53), (74022, 53), (296084,), (74022,))`

5.4. Perform data normalization/standardization

Data normalization/standardization is required so that features with higher scales do not dominate the model's performance. Hence all features should have same scale\

Data before scaling

```
In [49]: X_train.head()
```

```
Out[49]:
```

	loan_amnt	term	int_rate	annual_inc	dti	open_acc	revol_bal	revol_util	total_acc	any_mort	...	zip_c
133405	27000.0	1	16.29	82302.0	25.52	13.0	12014.0	48.6	29.0	1	...	
365868	6000.0	0	18.55	45000.0	19.37	8.0	3219.0	73.2	11.0	0	...	
71124	8975.0	0	9.71	65000.0	7.98	10.0	3932.0	34.5	58.0	0	...	
33923	9600.0	0	6.62	58000.0	25.01	10.0	57236.0	36.3	19.0	1	...	
30512	18000.0	0	11.53	75000.0	8.50	9.0	9916.0	35.8	17.0	0	...	

5 rows × 53 columns

```
In [50]: min_max_scaler = MinMaxScaler()
# Fit min_max_scaler to training data
min_max_scaler.fit(X_train)
# Scale the training and testing data
X_train = pd.DataFrame(min_max_scaler.transform(X_train), columns=X_train.columns)
X_test = pd.DataFrame(min_max_scaler.transform(X_test), columns=X_test.columns)
```

Data after scaling

```
In [51]: X_train.head()
```

```
Out[51]:
```

	loan_amnt	term	int_rate	annual_inc	dti	open_acc	revol_bal	revol_util	total_acc	any_mort	...	zip_c
0	0.697828	1.0	0.506230	0.308276	0.372229	0.48	0.154818	0.394481	0.457627	1.0	...	
1	0.144832	0.0	0.610521	0.161417	0.282526	0.28	0.041481	0.594156	0.152542	0.0	...	
2	0.223173	0.0	0.202584	0.240157	0.116394	0.36	0.050669	0.280032	0.949153	0.0	...	
3	0.239631	0.0	0.059991	0.212598	0.364790	0.36	0.737568	0.294643	0.288136	1.0	...	
4	0.460829	0.0	0.286571	0.279528	0.123979	0.32	0.127782	0.290584	0.254237	0.0	...	

5 rows × 53 columns

Check for imbalance in target class

```
In [52]: y_train.value_counts(normalize=True)*100
```

```
Out[52]:
loan_status
0      80.249186
1      19.750814
Name: proportion, dtype: float64
```

We can see a clear imbalance in the target class with **1** being ~**20%** and **0** being ~**80%**. Hence, I will use **SMOTE** to fix this imbalance

```
In [53]: sm = SMOTE(random_state=0)
```

```
X_train, y_train = sm.fit_resample(X_train, y_train)
y_train.value_counts(normalize=True)*100
```

```
Out[53]: loan_status
0      50.0
1      50.0
Name: proportion, dtype: float64
```

6. Build Logistic Regression model

Train the model


```
In [54]: model = LogisticRegression(solver='lbfgs', max_iter=300)
model.fit(X_train, y_train)
y_train_pred = model.predict(X_train)
```

Classification metrics and confusion matrix for Training data

```
In [55]: print(classification_report(y_train, y_train_pred))
```

	precision	recall	f1-score	support
0	0.81	0.80	0.80	237605
1	0.80	0.82	0.81	237605
accuracy			0.81	475210
macro avg	0.81	0.81	0.81	475210
weighted avg	0.81	0.81	0.81	475210

```
In [56]: # Confusion Matrix
cm = confusion_matrix(y_train, y_train_pred)
disp = ConfusionMatrixDisplay(cm)
disp.plot()
plt.title('Confusion Matrix')
plt.show()
```



Model prediction


```
In [57]: y_test_pred = model.predict(X_test)
```

Classification metrics and confusion matrix for Testing data

```
In [58]: print(classification_report(y_test, y_test_pred))
```

	precision	recall	f1-score	support
0	0.95	0.80	0.86	59619
1	0.49	0.81	0.61	14403
accuracy			0.80	74022
macro avg	0.72	0.80	0.74	74022
weighted avg	0.86	0.80	0.82	74022

```
In [59]: # Confusion Matrix
cm = confusion_matrix(y_test, y_test_pred)
disp = ConfusionMatrixDisplay(cm)
disp.plot()
plt.title('Confusion Matrix')
plt.show()
```




Insight

- Recall is high indicating that the model is able to identify 80% of the actual defaulters and 80% of non-defaulting customers
- Precision for class 1 (defaulters) is low. Of all the predicted defaulters, only 50% are actual defaulters
- With this model there is a risk of denying loans to deserving customers due to low precision score for defaulters

```
In [60]: feature_imp = pd.DataFrame({'Columns':X_train.columns, 'Coefficients':model.coef_[0]}).r

plt.figure(figsize=(8,8))
sns.barplot(data=feature_imp, y = 'Columns', x = 'Coefficients')
plt.title("Feature Importance for Model")
plt.yticks(fontsize=8)
plt.ylabel("Feature")
plt.tight_layout()
plt.show()
```



Insight

- The features **zip_code_29597**, **zip_code_05113**, **zip_code_00813**, **annual_inc** and **application_type_joint** have got high **positive weightage** and features **zip_code_86630**, **zip_code_11650**, **zip_code_93700**, **dti** and **open_acc** have got high **negative weightage** indicating their major contribution towards target variable

ROC and AUC


```
In [61]: # Predict probabilities for the test set
probs = model.predict_proba(X_test)[:,-1]

# Compute the false positive rate, true positive rate, and thresholds
fpr, tpr, thresholds = roc_curve(y_test, probs)

# Compute the area under the ROC curve
roc_auc = auc(fpr, tpr)
```



```
# Plot the ROC curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.show()
```



Insight

- ROC curve illustrates the trade off between TPR(True Positive Rate) and FPR(False Positive Rate) for various thresholds
- The AU-ROC value of 0.91 signifies that the model is able to differentiate well between the two classes
- Let us also look at PR Curve(Precision Recall Curve)


Precision Recall Curve

```
In [62]: precision, recall, thr = precision_recall_curve(y_test, probs)

# Area under Precision Recall Curve
apc = average_precision_score(y_test, probs)

# Plot the precision-recall curve
plt.plot(recall, precision, marker='.', label='PR curve (area = %0.2f)' % apc)
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
```

```
plt.legend(loc="lower left")
plt.show()
```




Insight

- PR curve illustrates the trade off between Precision and Recall for various thresholds
- The model has a AU-PRC value of 0.78 which is not that high. It is better than the random model which has a AU-PRC value of 0.5.
- This clearly indicates that we simply cannot conclude on the model's performance from just the ROC curve.

Find the threshold where precision and recall meet


```
In [63]: plt.figure()
plt.plot(thr,precision[0:len(thr)],label='precision',color='blue')
plt.plot(thr,recall[0:len(thr)],label='recall',color='orange')
intersection_thr = thr[np.where(precision == recall)[0][0]].round(4)
plt.axvline(intersection_thr, linestyle='--', color='red')
plt.text(intersection_thr, 0.01, str(intersection_thr), ha='left', color='red')
plt.title("Precision-recall curve at different thresholds")
plt.xlabel("Threshold values")
plt.ylabel("Precision and Recall values")
plt.legend(loc="upper right")
plt.grid()
plt.show()
```



```
In [64]: y_pred = model.predict_proba(X_test)[:,-1]
threshold_considered = intersection_thr
y_pred_custom = (y_pred > threshold_considered).astype('int')
print(classification_report(y_test, y_pred_custom))
```

	precision	recall	f1-score	support
0	0.92	0.92	0.92	59619
1	0.65	0.65	0.65	14403
accuracy			0.87	74022
macro avg	0.78	0.78	0.78	74022
weighted avg	0.87	0.87	0.87	74022

```
In [65]: # Confusion Matrix
cm = confusion_matrix(y_test, y_pred_custom)
disp = ConfusionMatrixDisplay(cm)
disp.plot()
plt.title('Confusion Matrix')
plt.show()
```



Insight

- With the new threshold, precision for class 1 (defaulters) has increased at a cost of decrease in Recall
- The overall F1-score and accuracy has increased

7. Insights

- **80%** of the customers have fully paid their loan and 20% are defaulters
- Loan amount and installment are highly **correlated** as it is obvious that high loan amount will have high installment amount
- Loan taken for short term, i.e. **3 years** are most likely to be **fully paid back**
- Most of the people have home ownership as **mortgage**
- Surprisingly, loans which are **not verified** are more likely to be **paid back**
- Loan taken as **joint** application type are more likely to be **paid back**
- People with **grade A** are more likely to fully pay their loan
- Loan taken for **wedding** are more likely to be paid back
- People from zip code **00813, 05113** fully pay back their loans whereas people from zip code **11650, 86630, 93700** are all defaulters
- The features **zip_code_29597, zip_code_05113, zip_code_00813, annual_inc, application_type_joint, zip_code_86630, zip_code_11650, zip_code_93700, dti, open_acc** affected the model outcome heavily
- As per the ROC curve and AU-ROC value of 0.91, the model is able to differentiate well between the defaulters and non-defaulters

- As per the PRC and AU-PRC value of 0.97, the model is able to return accurate results as well as return majority of all positive results(high recall)

8. Recommendation

- The bank can provide more short term loans, i.e. for 3 years, without much risk
- Provide more joint loans and scrutinize more individual and direct pay application types
- Analyze carefully the loan applications of customers with grades D, E, F and G. Do not provide them loans or provide smaller loans to these customers
- Reduce the loan given for small bussiness or analyze their application in detail before giving out loan to small bussiness
- Do not provide loans to customers with zip code 11650, 86630, 93700